# DATA_621_HW3

Chi Pong, Euclid Zhang, Jie Zou, Joseph Connolly, LeTicia Cancel

3/14/2022

```
train_df <- read.csv("https://raw.githubusercontent.com/ezaccountz/DATA_621/main/HW3/crime-training-dat
test_df <- read.csv("https://raw.githubusercontent.com/ezaccountz/DATA_621/main/HW3/crime-evaluation-da
```

# DATA EXPLORATION

## * Data Summary

```
summary(train_df)
```

```
##        zn             indus            chas              nox
##  Min.   :  0.00   Min.   : 0.460   Min.   :0.00000   Min.   :0.3890
##  1st Qu.:  0.00   1st Qu.: 5.145   1st Qu.:0.00000   1st Qu.:0.4480
##  Median :  0.00   Median : 9.690   Median :0.00000   Median :0.5380
##  Mean   : 11.58   Mean   :11.105   Mean   :0.07082   Mean   :0.5543
##  3rd Qu.: 16.25   3rd Qu.:18.100   3rd Qu.:0.00000   3rd Qu.:0.6240
##  Max.   :100.00   Max.   :27.740   Max.   :1.00000   Max.   :0.8710
##        rm             age             dis              rad
##  Min.   :3.863   Min.   :  2.90   Min.   : 1.130   Min.   : 1.00
##  1st Qu.:5.887   1st Qu.: 43.88   1st Qu.: 2.101   1st Qu.: 4.00
##  Median :6.210   Median : 77.15   Median : 3.191   Median : 5.00
##  Mean   :6.291   Mean   : 68.37   Mean   : 3.796   Mean   : 9.53
##  3rd Qu.:6.630   3rd Qu.: 94.10   3rd Qu.: 5.215   3rd Qu.:24.00
##  Max.   :8.780   Max.   :100.00   Max.   :12.127   Max.   :24.00
##       tax            ptratio          lstat             medv
##  Min.   :187.0   Min.   :12.6    Min.   : 1.730   Min.   : 5.00
##  1st Qu.:281.0   1st Qu.:16.9    1st Qu.: 7.043   1st Qu.:17.02
##  Median :334.5   Median :18.9    Median :11.350   Median :21.20
##  Mean   :409.5   Mean   :18.4    Mean   :12.631   Mean   :22.59
##  3rd Qu.:666.0   3rd Qu.:20.2    3rd Qu.:16.930   3rd Qu.:25.00
##  Max.   :711.0   Max.   :22.0    Max.   :37.970   Max.   :50.00
##      target
##  Min.   :0.0000
##  1st Qu.:0.0000
##  Median :0.0000
##  Mean   :0.4914
##  3rd Qu.:1.0000
##  Max.   :1.0000
```

From the summary:

- We see that there is **no missing value**
- All numeric variables, except **zn**, the max value is not significantly off from the third quartile. This indicates that they have no extream outliers. We will need to look at the distribution of **zn** to check the outliers.

## * Box Plots

```
data.m <- melt(train_df, id.vars = 'target')%>% mutate(target = as.factor(target))
ggplot(data.m, aes(x = variable, y = value, fill = target)) + geom_boxplot() +
  facet_wrap(~ variable, scales = 'free') + theme_classic()
```



From the boxplots, the following predictors seem to be good candidates differentiating target = 0 and target = 1: * **zn** * **indus** * **nox** * **age** * **dis** * **rad** * **tax**

Most of the predictors seem to have a different distribution for target = 0 and target = 1. Let's check the distributions of the predictors:

## * Distribution plots

```
target_factored <- as.factor(train_df$target)

plot_zn <- ggplot(train_df, aes(x=zn, color=target_factored)) + geom_density()
plot_indus <- ggplot(train_df, aes(x=indus, color=target_factored)) + geom_density()
plot_nox <- ggplot(train_df, aes(x=nox, color=target_factored)) + geom_density()
```

```r
plot_rm <- ggplot(train_df, aes(x=rm, color=target_factored)) + geom_density()
plot_age <- ggplot(train_df, aes(x=age, color=target_factored)) + geom_density()
plot_dis <- ggplot(train_df, aes(x=dis, color=target_factored)) + geom_density()
plot_rad <- ggplot(train_df, aes(x=rad, color=target_factored)) + geom_density()
plot_tax <- ggplot(train_df, aes(x=tax, color=target_factored)) + geom_density()
plot_prtatio <- ggplot(train_df, aes(x=ptratio, color=target_factored)) + geom_density()
plot_lstat <- ggplot(train_df, aes(x=lstat, color=target_factored)) + geom_density()
plots_medv <- ggplot(train_df, aes(x=medv, color=target_factored)) + geom_density()

plot_zn+plot_indus+plot_nox+plot_rm+plot_age+plot_dis+plot_rad+plot_tax+
  plot_prtatio+plot_lstat+plots_medv+plot_layout(ncol = 4, guides = "collect")
```
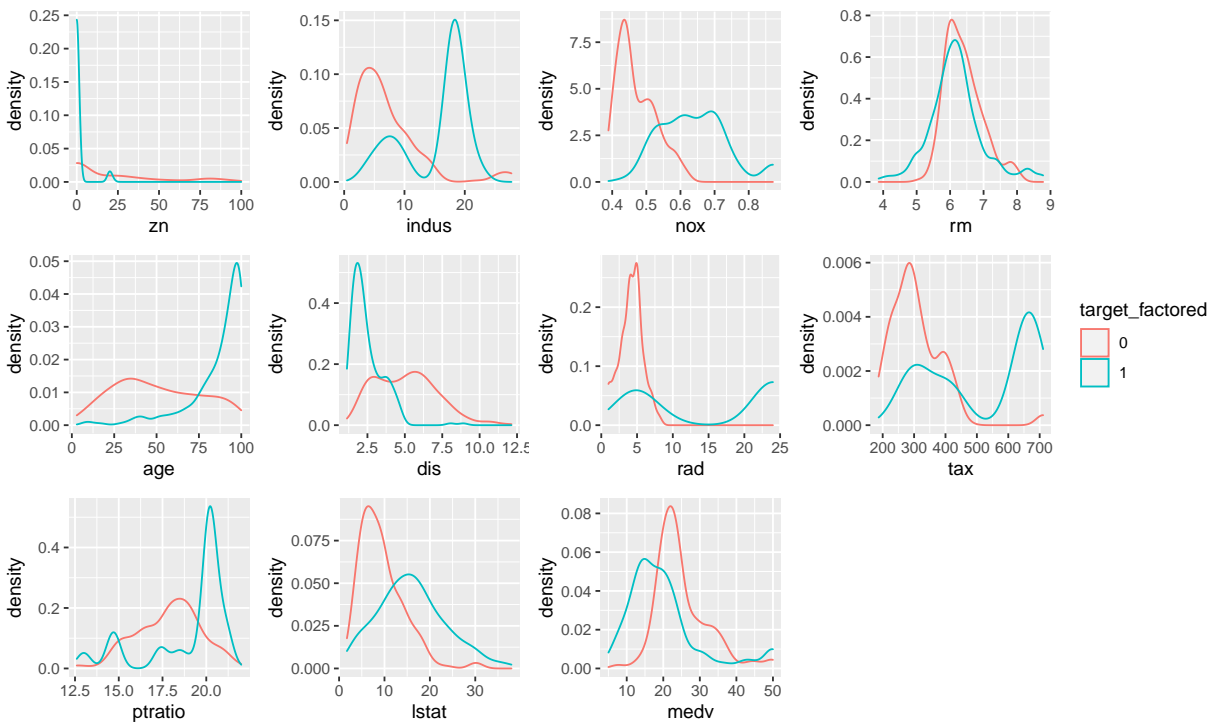
**zn** is zero-inflated for target = 1. We may want to add or transform it into a dummy variable indicating **zn** is greater than 0 or not. **lstat** and **medv** are right-skewed for both target = 0 and target = 1, we may consider a log-transformation. For other numeric variables, we may check later if transformations are needed.


## * Correlations

Now let's look at the correlations between the variables

```r
corrplot(cor(train_df, use = "na.or.complete"), method = 'number',
         type = 'lower', diag = FALSE, tl.srt = 0.1)
```

|       | zn    | indus | chas  | nox   | rm    | age   | dis   | rad   | tax   | ptratio | lstat | medv  |
|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|---------|-------|-------|
| indus | −0.54 |       |       |       |       |       |       |       |       |         |       |       |
| chas  | −0.04 | 0.06  |       |       |       |       |       |       |       |         |       |       |
| nox   | −0.52 | 0.76  | 0.10  |       |       |       |       |       |       |         |       |       |
| rm    | 0.32  | −0.39 | 0.09  | −0.30 |       |       |       |       |       |         |       |       |
| age   | −0.57 | 0.64  | 0.08  | 0.74  | −0.23 |       |       |       |       |         |       |       |
| dis   | 0.66  | −0.70 | −0.10 | −0.77 | 0.20  | −0.75 |       |       |       |         |       |       |
| rad   | −0.32 | 0.60  | −0.02 | 0.60  | −0.21 | 0.46  | −0.49 |       |       |         |       |       |
| tax   | −0.32 | 0.73  | −0.05 | 0.65  | −0.30 | 0.51  | −0.53 | 0.91  |       |         |       |       |
| ptratio | −0.39 | 0.39 | −0.13 | 0.18 | −0.36 | 0.26 | −0.23 | 0.47 | 0.47 |         |       |       |
| lstat | −0.43 | 0.61  | −0.05 | 0.60  | −0.63 | 0.61  | −0.51 | 0.50  | 0.56  | 0.38    |       |       |
| medv  | 0.38  | −0.50 | 0.16  | −0.43 | 0.71  | −0.38 | 0.26  | −0.40 | −0.49 | −0.52   | −0.74 |       |
| target | −0.43 | 0.60 | 0.08  | 0.73  | −0.15 | 0.63  | −0.62 | 0.63  | 0.61  | 0.25    | 0.47  | −0.27 |

We see that **rad** and **tax** have strong correlation.

Let's take a look at the linear plot of the two variables

```
plot(jitter(train_df$rad),jitter(train_df$tax), xlab="rad", ylab="tax")
abline(lm (train_df$tax ~ train_df$rad))
```

We can see that the correlation is strongly influenced by one point, where **rad** = 24.

Without **rad**, the correlation is only,

```
cor(train_df[train_df$rad < 20,"rad"],train_df[train_df$rad < 20,"tax"])
```

```
## [1] 0.1799913
```

In this case, we will not remove **rad** or **tax** from our model, but we will need to be cautious of the t-statistics of the two in our models.

## DATA PREPARATION

From the density plot of **zn**, we know that the variable is zero-inflated. The percentage of 0 values is

```
nrow(train_df[train_df$zn==0,])/nrow(train_df)
```

```
## [1] 0.7274678
```

Let's check the distribution of the **zn** without the 0 values

```
plot(density(train_df[train_df$zn>0,]$zn,na.rm=TRUE), main = "zn > 0")
```

**zn > 0**

Density

N = 127   Bandwidth = 8.982

The distribution looks a lot better.

We will add a new dummy variable zn_y indicating if **zn** is >0. The interaction **zn x zn_y = zn** so we don't need to do anything to it. If **zn_y** is deemed to be insignificant by our models, then we can simply drop it.

```
train_df$zn_y <- 0
train_df$zn_y[train_df$zn>0] <- 1
```

According to the text book *A Modern Approach To Regression With R*, "when the predictor variable X has a Poisson distribution, the log odds are a linear function of x". Let's check if any of the predictors follows a Poisson distribution

```
#Method of possion distribution test is from https://stackoverflow.com/questions/59809960/how-do-i-know

#two tail test
p_poisson <- function(x) {
  return (1-2 * abs((1 - pchisq((sum((x - mean(x))^2)/mean(x)), length(x) - 1))-0.5))
}

predictors <- colnames(train_df)
predictors <- predictors[!predictors %in% c("target","chas","zn_y")]

data.frame(mean_target0 = round(apply(train_df[train_df$target==0,predictors],2,mean),2),
    variance_target0 = round(apply(train_df[train_df$target==0,predictors],2,var),2),
    p_poisson_target0 = round(apply(train_df[train_df$target==0,predictors],2,p_poisson),2),
    mean_target1 = round(apply(train_df[train_df$target==1,predictors],2,mean),2),
```

```
    variance_target1 = round(apply(train_df[train_df$target==1,predictors],2,var),2),
    p_poisson_target1 = round(apply(train_df[train_df$target==1,predictors],2,p_poisson),2))
```

```
##          mean_target0 variance_target0 p_poisson_target0 mean_target1
## zn              21.48           850.75              0.00         1.33
## indus            7.04            30.27              0.00        15.31
## nox              0.47             0.00              0.00         0.64
## rm               6.40             0.31              0.00         6.18
## age             50.84           665.13              0.00        86.50
## dis              5.08             4.27              0.07         2.47
## rad              4.17             2.54              0.00        15.07
## tax            308.75          7956.24              0.00       513.77
## ptratio         17.86             3.35              0.00        18.96
## lstat            9.36            23.86              0.00        16.02
## medv            25.04            53.83              0.00        20.05
##          variance_target1 p_poisson_target1
## zn                  25.28                 0
## indus               29.28                 0
## nox                  0.01                 0
## rm                   0.66                 0
## age                297.90                 0
## dis                  1.16                 0
## rad                 90.53                 0
## tax              27786.70                 0
## ptratio              5.75                 0
## lstat               55.52                 0
## medv               105.65                 0
```

The null hypothesis in the tests is that the variable follows a poision distribution. Based on the p-values, we reject the null hypothesis for all predictors. None of the predictors follows a poisson distribution for both target = 0 and target = 1

Let's look at the distributions plots again.

```
plot_zn+plot_indus+plot_nox+plot_rm+plot_age+plot_dis+plot_rad+plot_tax+
  plot_prtatio+plot_lstat+plots_medv+plot_layout(ncol = 4, guides = "collect")
```
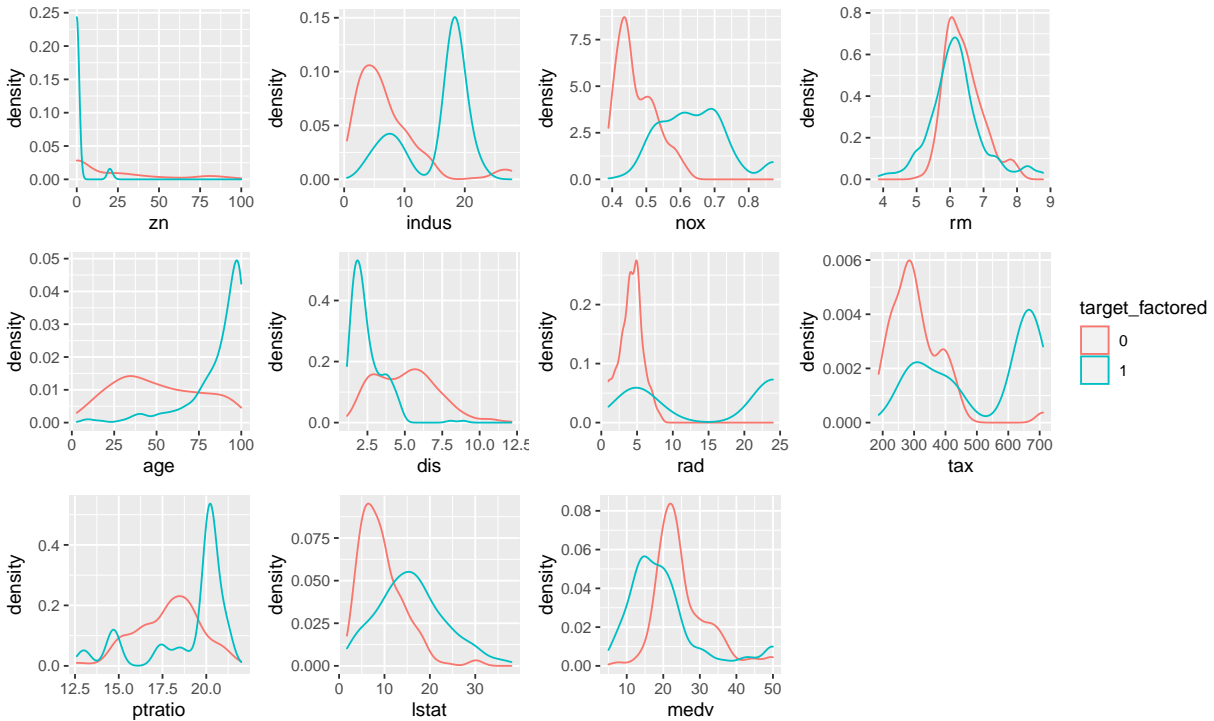
The distributions for rm with target = 0 and target = 1 are approximately normal with the same variance. Hence would not transform the variable The distributions for lstat and medv are skewed for both target = 0 and target = 1, we will add a log-transformed variable for each of them.

```
train_df$log_lstat <- log(train_df$lstat)
train_df$log_medv <- log(train_df$medv)
```

The distributions for indus, nox, age, dis, tax, ptratio look significantly different for the target values. Let perform a anova tests on the single predictor models to see if adding a log transformed or a quadratic transformed variable will improve the performance.

**Quadratic transformation test:**

```
predictors <- c("indus", "nox", "age", "dis", "tax", "ptratio")

n <- length(predictors)

model_compare <- data.frame(
    model_1 = paste0("target~",predictors),
    model_2 = paste0("target~",predictors,"+I(",predictors,"^2)"),
    Diff_DF = rep(0,n),
    Diff_Deviance = rep(0.0000,n),
    Pr_Gt_Chi = rep(0.0000,n)
)

for (i in (1:n)) {
    test_model_1 <- glm(target~train_df[,predictors[i]],family = binomial, train_df)
    test_model_2 <- glm(target~train_df[,predictors[i]]+
                        I(train_df[,predictors[i]]^2),family = binomial, train_df)
    anova_test <- anova(test_model_1,test_model_2,test="Chi")
```

```
    model_compare[i,3] <- anova_test$Df[2]
    model_compare[i,4] <- round(anova_test$Deviance[2],2)
    model_compare[i,5] <- round(anova_test$`Pr(>Chi)`[2],6)
}

model_compare
```

```
##            model_1                         model_2 Diff_DF Diff_Deviance Pr_Gt_Chi
## 1    target~indus    target~indus+I(indus^2)         1         31.13  0.000000
## 2      target~nox        target~nox+I(nox^2)         1          0.29  0.587879
## 3      target~age        target~age+I(age^2)         1          7.63  0.005748
## 4      target~dis        target~dis+I(dis^2)         1          3.64  0.056401
## 5      target~tax        target~tax+I(tax^2)         1          0.44  0.505781
## 6 target~ptratio target~ptratio+I(ptratio^2)         1         98.71  0.000000
```

**Log transformation test:**

```
predictors <- c("indus", "nox", "age", "dis", "tax", "ptratio")

n <- length(predictors)

model_compare <- data.frame(
    model_1 = paste0("target~",predictors),
    model_2 = paste0("target~",predictors,"+I(log(",predictors,"))"),
    Diff_DF = rep(0,n),
    Diff_Deviance = rep(0.0000,n),
    Pr_Gt_Chi = rep(0.0000,n)
)

for (i in (1:n)) {
    test_model_1 <- glm(target~train_df[,predictors[i]],family = binomial, train_df)
    test_model_2 <- glm(target~train_df[,predictors[i]]+
                    I(log(train_df[,predictors[i]])),family = binomial, train_df)
    anova_test <- anova(test_model_1,test_model_2,test="Chi")
    model_compare[i,3] <- anova_test$Df[2]
    model_compare[i,4] <- round(anova_test$Deviance[2],2)
    model_compare[i,5] <- round(anova_test$`Pr(>Chi)`[2],6)
}

model_compare
```

```
##            model_1                          model_2 Diff_DF Diff_Deviance Pr_Gt_Chi
## 1    target~indus    target~indus+I(log(indus))         1         13.91  0.000192
## 2      target~nox        target~nox+I(log(nox))         1          0.63  0.427570
## 3      target~age        target~age+I(log(age))         1          6.37  0.011603
## 4      target~dis        target~dis+I(log(dis))         1          5.15  0.023182
## 5      target~tax        target~tax+I(log(tax))         1          1.00  0.317895
## 6 target~ptratio target~ptratio+I(log(ptratio))         1         98.09  0.000000
```

For indus, the improvement is bigger by adding the squared term. For ptratio, since the distribution is left-skewed, it may be better to add the squared term. For other variables, no transformation is added.

```
train_df$indus_squared <- train_df$indus^2
train_df$ptratio_squared <- train_df$ptratio^2
```

**Interaction term test:**

**chas** is a dummy variable. We will perform a anova tests on the single predictor models to see if adding an interaction between **chas** and a predictor will improve the model.

```
predictors <- colnames(train_df)
predictors <- predictors[!predictors %in% c("target","chas","zn_y","log_lstat",
                                            "log_medv","indus_squared",
                                            "ptratio_squared")]

interaction_test <- data.frame(matrix(ncol = 3, nrow = 0))
colnames(interaction_test) <- c("Preditor","Interaction","Pr_Gt_Chi")
class(interaction_test$Pr_Gt_Chi) = "Numeric"

for (predictor in predictors) {
  interaction_test[nrow(interaction_test) + 1,] <-
    c(predictor, paste0(predictor, ":chas"),
      round(anova(glm(target ~ train_df[,predictor]*chas,data = train_df,
                      family = "binomial"),test="Chi")[4,5],4))
}
```

```
interaction_test
```

```
##      Preditor  Interaction Pr_Gt_Chi
## 1          zn      zn:chas    0.0645
## 2       indus   indus:chas     0.954
## 3         nox     nox:chas    0.6638
## 4          rm      rm:chas    0.6647
## 5         age     age:chas    0.0719
## 6         dis     dis:chas    0.0681
## 7         rad     rad:chas    0.0191
## 8         tax     tax:chas         0
## 9     ptratio ptratio:chas    0.3917
## 10      lstat   lstat:chas    0.1006
## 11       medv    medv:chas    0.1555
```

From the result, we will add an interaction between **tax** and **chas** and an interaction between **rad** and **chas** to our preditor candidates.

```
train_df$tax_chas <- train_df$tax * train_df$chas
train_df$rad_chas <- train_df$rad * train_df$chas
```

# BUILD MODELS

## * 1. Full model:

The full model includes all original and the transformed verson of the predictors

```
full_model <- glm(target~.,family = binomial, train_df)

#store the model formulas for buidling models for cross validation
model_formulas <- c(paste(deparse(formula(full_model), width.cutoff = 500), collapse=""))
```

```
summary(full_model)
```

```
##
## Call:
## glm(formula = target ~ ., family = binomial, data = train_df)
##
## Deviance Residuals:
##     Min       1Q   Median       3Q      Max
## -2.6227  -0.0932   0.0000   0.0001   3.7950
##
## Coefficients:
##                  Estimate Std. Error z value Pr(>|z|)
## (Intercept)     5.352e+01  2.666e+01   2.008 0.044688 *
## zn              1.010e-02  3.979e-02   0.254 0.799658
## indus           1.342e+00  4.258e-01   3.152 0.001623 **
## chas           -7.255e+02  5.172e+04  -0.014 0.988806
## nox             3.688e+01  8.993e+00   4.100 4.13e-05 ***
## rm             -1.801e+00  9.893e-01  -1.820 0.068699 .
## age             3.642e-02  1.562e-02   2.332 0.019712 *
## dis             6.235e-01  2.946e-01   2.116 0.034323 *
## rad             1.219e+00  3.188e-01   3.824 0.000131 ***
## tax            -2.145e-02  7.854e-03  -2.731 0.006311 **
## ptratio        -6.900e+00  2.303e+00  -2.996 0.002732 **
## lstat           2.391e-01  1.681e-01   1.422 0.154981
## medv            5.679e-01  2.286e-01   2.484 0.012980 *
## zn_y           -2.085e+00  1.398e+00  -1.491 0.136052
## log_lstat      -3.186e+00  2.269e+00  -1.404 0.160181
## log_medv       -8.051e+00  4.937e+00  -1.631 0.102965
## indus_squared  -4.242e-02  1.329e-02  -3.191 0.001418 **
## ptratio_squared 2.027e-01  6.390e-02   3.172 0.001516 **
## tax_chas        2.868e+00  2.052e+02   0.014 0.988849
## rad_chas       -1.707e+01  1.429e+03  -0.012 0.990469
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##     Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 148.14  on 446  degrees of freedom
## AIC: 188.14
##
## Number of Fisher Scoring iterations: 22
```

## * 2. Backward Elimination by AIC:

Starting with our full model, perform backward elimination by comparing the AIC of the models.

```
model_AIC <- step(full_model, trace=0)
model_formulas <- c(model_formulas, paste(deparse(formula(model_AIC)),
                                            width.cutoff = 500), collapse=""))
```

The result model is:

```
summary(model_AIC)
```

```
##
## Call:
## glm(formula = target ~ indus + chas + nox + rm + age + dis +
##     rad + tax + ptratio + lstat + medv + zn_y + log_lstat + log_medv +
##     indus_squared + ptratio_squared + tax_chas, family = binomial,
##     data = train_df)
##
## Deviance Residuals:
##     Min       1Q   Median       3Q      Max
## -2.6018  -0.0948   0.0000   0.0001   3.7829
##
## Coefficients:
##                   Estimate Std. Error z value Pr(>|z|)
## (Intercept)      5.453e+01  2.645e+01   2.061 0.039262 *
## indus            1.325e+00  4.199e-01   3.156 0.001601 **
## chas            -5.397e+03  3.651e+05  -0.015 0.988205
## nox              3.637e+01  8.744e+00   4.160 3.18e-05 ***
## rm              -1.813e+00  9.855e-01  -1.840 0.065813 .
## age              3.602e-02  1.550e-02   2.324 0.020109 *
## dis              6.136e-01  2.945e-01   2.083 0.037218 *
## rad              1.210e+00  3.177e-01   3.810 0.000139 ***
## tax             -2.117e-02  7.793e-03  -2.716 0.006603 **
## ptratio         -6.908e+00  2.298e+00  -3.006 0.002649 **
## lstat            2.396e-01  1.684e-01   1.423 0.154768
## medv             5.748e-01  2.267e-01   2.535 0.011244 *
## zn_y            -1.864e+00  1.097e+00  -1.699 0.089327 .
## log_lstat       -3.227e+00  2.267e+00  -1.424 0.154519
## log_medv        -8.226e+00  4.900e+00  -1.679 0.093179 .
## indus_squared   -4.183e-02  1.306e-02  -3.202 0.001365 **
## ptratio_squared  2.029e-01  6.378e-02   3.181 0.001468 **
## tax_chas         1.949e+01  1.318e+03   0.015 0.988204
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##     Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 148.20  on 448  degrees of freedom
## AIC: 184.2
##
## Number of Fisher Scoring iterations: 25
```

## * 3. Backward Elimination with Chi-square test:

Starting with our full model, perform backward elimination with Chi-square test.

```r
#Define a function to perform backward elimination with Chi-square test
#using the significancy / alpha as one of the parameters

backward_chi <- function (train_df, significancy) {
  glm_string <- "target~."
  glm_formula <- as.formula(glm_string)

  repeat{
    drop1_chi <- drop1(glm(glm_formula, family=binomial, train_df), test="Chi")

    chi_result <- data.frame(preditors = rownames(drop1_chi)[-1],
            p_value = drop1_chi[-1,5])
    chi_result <- chi_result[order(chi_result$p_value,decreasing=TRUE),]

    if(chi_result[1,2] < significancy){
        break
    }
    else {
        glm_string <- paste0(glm_string,"-",chi_result[1,1])
        glm_formula <- as.formula(glm_string)
    }
  }

  return(glm_formula)
}
```

**model with alpha 0.1 (based on Chi-square test)**

```r
model_chi_0.1 <- backward_chi(train_df, 0.1)
model_formulas <- c(model_formulas, model_chi_0.1)
model_chi_0.1 <- glm(model_chi_0.1, family=binomial, train_df)

summary(model_chi_0.1)
```

```
##
## Call:
## glm(formula = model_chi_0.1, family = binomial, data = train_df)
##
## Deviance Residuals:
##     Min       1Q   Median       3Q      Max
## -2.7406  -0.0828   0.0000   0.0001   3.7127
##
## Coefficients:
##                  Estimate Std. Error z value Pr(>|z|)
## (Intercept)     5.201e+01  2.542e+01   2.046 0.040793 *
## indus           1.359e+00  4.286e-01   3.171 0.001521 **
## chas           -5.510e+03  3.890e+05  -0.014 0.988699
## nox             3.631e+01  8.729e+00   4.160 3.18e-05 ***
## rm             -1.622e+00  8.247e-01  -1.966 0.049273 *
## age             3.496e-02  1.462e-02   2.391 0.016786 *
## dis             6.704e-01  2.935e-01   2.284 0.022385 *
## rad             1.224e+00  3.178e-01   3.851 0.000118 ***
## tax            -2.186e-02  7.821e-03  -2.795 0.005195 **
```

```
## ptratio        -6.725e+00  2.261e+00  -2.974 0.002935 **
## medv            7.044e-01  2.064e-01   3.412 0.000645 ***
## zn_y           -1.858e+00  1.078e+00  -1.724 0.084774 .
## log_medv       -1.090e+01  4.489e+00  -2.428 0.015167 *
## indus_squared  -4.276e-02  1.335e-02  -3.203 0.001358 **
## ptratio_squared 1.978e-01  6.285e-02   3.148 0.001643 **
## tax_chas        1.989e+01  1.404e+03   0.014 0.988698
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##     Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 150.31  on 450  degrees of freedom
## AIC: 182.31
##
## Number of Fisher Scoring iterations: 25
```

**model with alpha 0.05 (based on Chi-square test)**

```
model_chi_0.05 <- backward_chi(train_df, 0.05)
model_formulas <- c(model_formulas, model_chi_0.05)
model_chi_0.05 <- glm(model_chi_0.05, family=binomial, train_df)
summary(model_chi_0.05)
```

```
##
## Call:
## glm(formula = model_chi_0.05, family = binomial, data = train_df)
##
## Deviance Residuals:
##      Min       1Q    Median       3Q       Max
## -2.36990  -0.09297   0.00000   0.00003   3.05262
##
## Coefficients:
##                  Estimate Std. Error z value Pr(>|z|)
## (Intercept)      6.020e+01  2.251e+01   2.674 0.007490 **
## indus            1.407e+00  3.722e-01   3.780 0.000157 ***
## chas            -5.411e+03  2.712e+05  -0.020 0.984082
## nox              3.217e+01  6.979e+00   4.610 4.02e-06 ***
## rm              -6.657e-01  6.996e-01  -0.952 0.341302
## rad              1.311e+00  2.903e-01   4.517 6.26e-06 ***
## tax             -2.429e-02  6.625e-03  -3.667 0.000246 ***
## ptratio         -6.620e+00  2.173e+00  -3.046 0.002319 **
## medv             6.507e-01  2.025e-01   3.212 0.001316 **
## log_medv        -1.315e+01  4.484e+00  -2.933 0.003358 **
## indus_squared   -4.418e-02  1.160e-02  -3.808 0.000140 ***
## ptratio_squared  1.950e-01  6.093e-02   3.200 0.001375 **
## tax_chas         1.954e+01  9.790e+02   0.020 0.984078
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
```

```
##     Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 160.19  on 453  degrees of freedom
## AIC: 186.19
##
## Number of Fisher Scoring iterations: 24
```

**model with alpha 0.001 (based on Chi-square test)**

```
model_chi_0.001 <- backward_chi(train_df, 0.001)
model_formulas <- c(model_formulas, model_chi_0.001)
model_chi_0.001 <- glm(model_chi_0.001, family=binomial, train_df)
summary(model_chi_0.001)
```

```
##
## Call:
## glm(formula = model_chi_0.001, family = binomial, data = train_df)
##
## Deviance Residuals:
##      Min        1Q    Median        3Q       Max
## -2.30930  -0.10815   0.00000   0.00013   2.94594
##
## Coefficients:
##                    Estimate Std. Error z value Pr(>|z|)
## (Intercept)       3.571e+01  2.051e+01   1.741 0.081663 .
## indus             1.188e+00  3.649e-01   3.256 0.001132 **
## chas             -5.464e+03  2.914e+05  -0.019 0.985041
## nox               3.285e+01  7.007e+00   4.688 2.76e-06 ***
## rm               -1.518e-01  6.562e-01  -0.231 0.817100
## rad               1.165e+00  2.953e-01   3.947 7.92e-05 ***
## tax              -2.195e-02  6.597e-03  -3.327 0.000877 ***
## ptratio          -7.218e+00  2.213e+00  -3.262 0.001107 **
## medv              8.703e-02  6.717e-02   1.295 0.195148
## indus_squared    -3.794e-02  1.132e-02  -3.353 0.000801 ***
## ptratio_squared   2.110e-01  6.166e-02   3.421 0.000623 ***
## tax_chas          1.973e+01  1.052e+03   0.019 0.985036
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##     Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 169.78  on 454  degrees of freedom
## AIC: 193.78
##
## Number of Fisher Scoring iterations: 24
```

## * 4. Backward Elimination based on the t-values of the coefficients:

Starting with our full model, perform backward elimination based on the t-values of the coefficients.

```
#Define a function to perform backward elimination based on the t-values of the coefficients
#using the significancy / alpha as one of the parameters
```

```r
backward_p <- function (train_df, significancy) {
  glm_string <- "target~."
  glm_formula <- as.formula(glm_string)

  repeat{
    model_p <- glm(glm_formula, family=binomial, train_df)

    p_result <- data.frame(preditors = rownames(summary(model_p)$coefficients)[-1],
             p_value = summary(model_p)$coefficients[-1,4])
    p_result <- p_result[order(p_result$p_value,decreasing=TRUE),]


    if(p_result[1,2] < significancy){
        break
    }
    else {
        glm_string <- paste0(glm_string,"-",p_result[1,1])
        glm_formula <- as.formula(glm_string)
    }
  }

  return(glm_formula)
}
```

**model with alpha 0.05 (based on the t-values of the coefficients)**

alpha = 0.1 produces the same model as alpha = 0.05 so alpha = 0.1 is not used here.

```r
model_p_0.05 <- backward_p(train_df, 0.05)
model_formulas <- c(model_formulas, model_p_0.05)
model_p_0.05 <- glm(model_p_0.05, family=binomial, train_df)
summary(model_p_0.05)
```

```
##
## Call:
## glm(formula = model_p_0.05, family = binomial, data = train_df)
##
## Deviance Residuals:
##     Min      1Q   Median      3Q      Max
## -2.1811  -0.1348  -0.0043   0.0012   3.7040
##
## Coefficients:
##                 Estimate Std. Error z value Pr(>|z|)
## (Intercept)    21.565494  21.025464   1.026 0.305040
## indus           0.655387   0.270470   2.423 0.015387 *
## nox            32.968725   8.347166   3.950 7.83e-05 ***
## age             0.025691   0.011714   2.193 0.028298 *
## dis             0.666502   0.266970   2.497 0.012541 *
## rad             0.900580   0.211886   4.250 2.13e-05 ***
## tax            -0.013018   0.005438  -2.394 0.016672 *
## ptratio        -6.272631   2.063883  -3.039 0.002372 **
## lstat           0.100858   0.049742   2.028 0.042600 *
## medv            0.176336   0.053238   3.312 0.000926 ***
```

16

```
## zn_y             -2.617454   1.063343  -2.462 0.013834 *
## indus_squared    -0.021730   0.008692  -2.500 0.012414 *
## ptratio_squared   0.182663   0.056883   3.211 0.001322 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##     Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 172.68  on 453  degrees of freedom
## AIC: 198.68
##
## Number of Fisher Scoring iterations: 9
```

**model with alpha 0.01 (based on the t-values of the coefficients)**

```
model_p_0.01 <- backward_p(train_df, 0.01)
model_formulas <- c(model_formulas, model_p_0.01)
model_p_0.01 <- glm(model_p_0.01, family=binomial, train_df)
summary(model_p_0.01)
```

```
##
## Call:
## glm(formula = model_p_0.01, family = binomial, data = train_df)
##
## Deviance Residuals:
##      Min        1Q    Median        3Q       Max
## -1.88094  -0.18124  -0.00327   0.00014   2.85404
##
## Coefficients:
##                  Estimate Std. Error z value Pr(>|z|)
## (Intercept)     34.400490  17.446746   1.972 0.048639 *
## indus            0.831323   0.282877   2.939 0.003295 **
## nox             29.226546   6.347383   4.605 4.13e-06 ***
## rad              1.132391   0.224553   5.043 4.59e-07 ***
## tax             -0.017648   0.005160  -3.420 0.000626 ***
## ptratio         -6.515603   1.893531  -3.441 0.000580 ***
## indus_squared   -0.026394   0.008849  -2.983 0.002857 **
## ptratio_squared  0.188472   0.052980   3.557 0.000375 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##     Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 193.89  on 458  degrees of freedom
## AIC: 209.89
##
## Number of Fisher Scoring iterations: 9
```

**model with alpha 0.001 (based on the t-values of the coefficients)**

```r
model_p_0.001 <- backward_p(train_df, 0.001)
model_formulas <- c(model_formulas, model_p_0.001)
model_p_0.001 <- glm(model_p_0.001, family=binomial, train_df)
summary(model_p_0.001)
```

```
##
## Call:
## glm(formula = model_p_0.001, family = binomial, data = train_df)
##
## Deviance Residuals:
##      Min       1Q    Median       3Q      Max
## -1.89721  -0.27798  -0.03997   0.00557   2.55954
##
## Coefficients:
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept) -19.867422   2.368325  -8.389  < 2e-16 ***
## nox          35.633515   4.523677   7.877 3.35e-15 ***
## rad           0.637643   0.119444   5.338 9.38e-08 ***
## tax          -0.008146   0.002332  -3.493 0.000478 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##     Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 224.47  on 462  degrees of freedom
## AIC: 232.47
##
## Number of Fisher Scoring iterations: 8
```

## SELECT MODELS

First, let's compare different metrics of all models we have built

```r
models <- list(full_model, model_AIC, model_chi_0.1, model_chi_0.05, model_chi_0.001,
               model_p_0.05, model_p_0.01, model_p_0.001)
model_names <- list("full_model", "model_AIC", "model_chi_0.1", "model_chi_0.05",
                    "model_chi_0.001","model_p_0.05", "model_p_0.01", "model_p_0.001")

model_compare <- data.frame(
    model = rep("",length(models)),
    Deviance = rep(0.0000,length(models)),
    AIC = rep(0.0000,length(models)),
    Accuracy = rep(0.0000,length(models)),
    Sensitivity = rep(0.0000,length(models)),
    Specificity = rep(0.0000,length(models)),
    Precision = rep(0.0000,length(models)),
    F1 = rep(0.0000,length(models)),
    AUC = rep(0.0000,length(models)),
    Nagelkerke_R_squared = rep(0.0000,length(models))
)
```

```r
for (i in c(1:length(models))) {
  predicted_class <- ifelse(models[[i]]$fitted.values>0.5,1,0)
  confusion_matrix <- confusionMatrix(as.factor(predicted_class),
                                      as.factor(train_df$target),positive = "1")

  model_compare[i,1] <- model_names[i]
  model_compare[i,2] <- round(models[[i]]$deviance,4)
  model_compare[i,3] <- models[[i]]$aic
  model_compare[i,4] <- confusion_matrix$overall[1]
  model_compare[i,5] <- confusion_matrix$byClass[1]
  model_compare[i,6] <- confusion_matrix$byClass[2]
  model_compare[i,7] <- confusion_matrix$byClass[3]
  model_compare[i,8] <- 2*confusion_matrix$byClass[1]*confusion_matrix$byClass[3]/
                        (confusion_matrix$byClass[1]+confusion_matrix$byClass[3])
  model_compare[i,9] <- auc(roc(train_df$target, models[[i]]$fitted.values))
  model_compare[i,10] <- (1-exp((models[[i]]$dev-models[[i]]$null)/
                                length(models[[i]]$residuals)))/
                         (1-exp(-models[[i]]$null/length(models[[i]]$residuals)))

}
model_compare
```

```
##                 model Deviance      AIC  Accuracy Sensitivity Specificity Precision
## 1         full_model 148.1421 188.1421 0.9356223   0.9257642   0.9451477 0.9422222
## 2          model_AIC 148.1999 184.1999 0.9356223   0.9257642   0.9451477 0.9422222
## 3    model_chi_0.1 150.3125 182.3125 0.9442060   0.9301310   0.9578059 0.9551570
## 4   model_chi_0.05 160.1878 186.1878 0.9420601   0.9257642   0.9578059 0.9549550
## 5 model_chi_0.001 169.7788 193.7788 0.9291845   0.9213974   0.9367089 0.9336283
## 6      model_p_0.05 172.6756 198.6756 0.9227468   0.9170306   0.9282700 0.9251101
## 7      model_p_0.01 193.8917 209.8917 0.9163090   0.9213974   0.9113924 0.9094828
## 8     model_p_0.001 224.4719 232.4719 0.8884120   0.8384279   0.9367089 0.9275362
##          F1       AUC Nagelkerke_R_squared
## 1 0.9339207 0.9854808            0.8752040
## 2 0.9339207 0.9856466            0.8751471
## 3 0.9424779 0.9844306            0.8730647
## 4 0.9401330 0.9814457            0.8632040
## 5 0.9274725 0.9794926            0.8534249
## 6 0.9210526 0.9789030            0.8504315
## 7 0.9154013 0.9747572            0.8279318
## 8 0.8807339 0.9594458            0.7936446
```

Since this is **logistic regression with binary data**, Deviance shouldn't be used to judge a model's goodness of fit. **We will mainly use AIC and the accuracy**. Depending on the business objective, we may use other metrics such as sensitivity and specificity to compare the models' performance. However, the business objective is not defined here so we simply use the accuracy. The Nagelkerke R squared is a pseudo version of the R squared, since R squared can not be used for generalized linear regression. The Nagelkerke R squared should not be used to judge the goodness of fit of a single model. It can be used to compare the fit of different models.

The result shows that model_chi_0.1 has the lowest AIC and best performance in predicting the using the training data. The models produced based the t-values of the coefficients are not doing so well. It is reasonable since some of the predictors have high correlation with each other.

The following is the confusion matrix for our best model model_chi_0.1:

```
predicted_class <- ifelse(model_chi_0.1$fitted.values>0.5,1,0)
confusion_matrix <- confusionMatrix(as.factor(predicted_class),
                                    as.factor(train_df$target),positive = "1")
confusion_matrix
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   0   1
##          0 227  16
##          1  10 213
##
##                Accuracy : 0.9442
##                  95% CI : (0.9193, 0.9632)
##     No Information Rate : 0.5086
##     P-Value [Acc > NIR] : <2e-16
##
##                   Kappa : 0.8883
##
##  Mcnemar's Test P-Value : 0.3268
##
##             Sensitivity : 0.9301
##             Specificity : 0.9578
##          Pos Pred Value : 0.9552
##          Neg Pred Value : 0.9342
##              Prevalence : 0.4914
##          Detection Rate : 0.4571
##    Detection Prevalence : 0.4785
##       Balanced Accuracy : 0.9440
##
##        'Positive' Class : 1
##
```

## * Cross Validation (5 fold)

Let's perform a cross validation on all the models we have to check if they are doing well with unseen data

```
models <- list(full_model, model_AIC, model_chi_0.1, model_chi_0.05, model_chi_0.001,
               model_p_0.05, model_p_0.01, model_p_0.001)
model_names <- list("full_model", "model_AIC", "model_chi_0.1", "model_chi_0.05",
                    "model_chi_0.001","model_p_0.05", "model_p_0.01", "model_p_0.001")


model_compare <- data.frame(
    model = rep("",length(models)),
    Accuracy_1 = rep(0.0000,length(models)),
    Accuracy_2 = rep(0.0000,length(models)),
    Accuracy_3 = rep(0.0000,length(models)),
    Accuracy_4 = rep(0.0000,length(models)),
    Accuracy_5 = rep(0.0000,length(models)),
```

```
    Accuracy_average = rep(0.0000,length(models)),
    AIC_1 = rep(0.0000,length(models)),
    AIC_2 = rep(0.0000,length(models)),
    AIC_3 = rep(0.0000,length(models)),
    AIC_4 = rep(0.0000,length(models)),
    AIC_5 = rep(0.0000,length(models)),
    AIC_average = rep(0.0000,length(models))
)
```

```
set.seed(14)
cv_df<-train_df[sample(nrow(train_df)),]
folds <- cut(seq(1,nrow(cv_df)),breaks=5,labels=FALSE)

#Perform 5 fold cross validation
for(i in 1:5){

    testIndexes <- which(folds==i,arr.ind=TRUE)
    testData <- cv_df[testIndexes, ]
    trainData <- cv_df[-testIndexes, ]

    for (j in c(1:length(models))) {
        test_model <- glm(model_formulas[[j]], family=binomial, trainData)
        predicted_class <- ifelse(predict(test_model,testData,type="response")>0.5,1,0)
        confusion_matrix <- confusionMatrix(as.factor(predicted_class),
                                as.factor(testData$target),positive = "1")
        model_compare[j,1+i] <- confusion_matrix$overall[1]
        model_compare[j,7+i] <- test_model$aic
    }


}

model_compare$model <- unlist(model_names)
model_compare$Accuracy_average <- apply(model_compare[,c(2:6)],1,mean)
model_compare$AIC_average <- apply(model_compare[,c(8:12)],1,mean)
```

The following table shows the accuracy of predictions with the test data and the AICs of the trained model.

```
model_compare
```

```
##              model Accuracy_1 Accuracy_2 Accuracy_3 Accuracy_4 Accuracy_5
## 1      full_model  0.9042553  0.8924731  0.9247312  0.9462366  0.9139785
## 2       model_AIC  0.9042553  0.8924731  0.9247312  0.9462366  0.9139785
## 3   model_chi_0.1  0.9042553  0.8817204  0.9247312  0.9462366  0.9139785
## 4  model_chi_0.05  0.8936170  0.9247312  0.9462366  0.9354839  0.9139785
## 5 model_chi_0.001  0.8404255  0.9032258  0.9247312  0.9354839  0.8924731
## 6     model_p_0.05  0.8723404  0.8924731  0.9139785  0.9354839  0.8924731
## 7     model_p_0.01  0.8936170  0.8817204  0.9032258  0.9462366  0.8817204
## 8    model_p_0.001  0.7978723  0.9032258  0.9032258  0.9032258  0.8494624
##   Accuracy_average    AIC_1    AIC_2    AIC_3    AIC_4    AIC_5 AIC_average
## 1        0.9163349 138.2636 153.1110 154.1078 165.8600 162.2430    154.7171
## 2        0.9163349 134.2647 149.1812 150.1884 161.9548 158.3747    150.7928
## 3        0.9141844 131.6432 151.1524 146.5266 160.1229 156.9649    149.2820
```

21

```
## 4         0.9228094 136.2853 150.3605 153.8125 159.8476 160.0786      152.0769
## 5         0.8992679 141.4014 157.6030 159.1907 165.3722 164.1969      157.5529
## 6         0.9013498 147.3185 158.6745 157.4345 172.7835 166.4235      160.5269
## 7         0.9013040 156.1461 170.6614 169.5861 178.6296 172.1746      169.4396
## 8         0.8714024 168.8905 192.8978 188.4387 193.8342 190.5635      186.9249
```

model_AIC and model_chi_0.1 have the best performance. Since model_chi_0.1 is a simpler model with 15 coefficients, we select model_chi_0.1 to be our best model as it is a more *parsimonious* model

```
length(model_chi_0.1$coefficients) - 1  # -1 for the intercept
```

```
## [1] 15
```

```
length(model_AIC$coefficients) - 1     # -1 for the intercept
```

```
## [1] 17
```

Let's check our final model again.

```
summary(model_chi_0.1)
```

```
##
## Call:
## glm(formula = model_chi_0.1, family = binomial, data = train_df)
##
## Deviance Residuals:
##     Min       1Q   Median       3Q      Max
## -2.7406  -0.0828   0.0000   0.0001   3.7127
##
## Coefficients:
##                   Estimate Std. Error z value Pr(>|z|)
## (Intercept)      5.201e+01  2.542e+01   2.046 0.040793 *
## indus            1.359e+00  4.286e-01   3.171 0.001521 **
## chas            -5.510e+03  3.890e+05  -0.014 0.988699
## nox              3.631e+01  8.729e+00   4.160 3.18e-05 ***
## rm              -1.622e+00  8.247e-01  -1.966 0.049273 *
## age              3.496e-02  1.462e-02   2.391 0.016786 *
## dis              6.704e-01  2.935e-01   2.284 0.022385 *
## rad              1.224e+00  3.178e-01   3.851 0.000118 ***
## tax             -2.186e-02  7.821e-03  -2.795 0.005195 **
## ptratio         -6.725e+00  2.261e+00  -2.974 0.002935 **
## medv             7.044e-01  2.064e-01   3.412 0.000645 ***
## zn_y            -1.858e+00  1.078e+00  -1.724 0.084774 .
## log_medv        -1.090e+01  4.489e+00  -2.428 0.015167 *
## indus_squared   -4.276e-02  1.335e-02  -3.203 0.001358 **
## ptratio_squared  1.978e-01  6.285e-02   3.148 0.001643 **
## tax_chas         1.989e+01  1.404e+03   0.014 0.988698
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
```

```
##
##     Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 150.31  on 450  degrees of freedom
## AIC: 182.31
##
## Number of Fisher Scoring iterations: 25
```

*zn_y* is not so significant. However, from the distribution plots, **zn** is has strong ability to differentiate target = 0 and target = 1 when **zn** is 0. It does poorly when **zn** is 1. We should keep this in our model.

For **chas** and **tax_chas**, they are highly correlated since **tax_chas** = 0 when **chas** = 0. The percentage of 0 in **chas** is

```
nrow(train_df[train_df$chas == 0,])/nrow(train_df)
```

```
## [1] 0.9291845
```

Since they are correlated, we would not judge the two coefficients by the t-value. The Chi-square tests told us that these two variables are important to the model's performance. We will keep **chas** and **tax_chas** in our model.

## * model diagnostics

Now let's look at the marginal plots to see if our model is fitting well to the training data

```
marginalModelPlots(model_chi_0.1,~nox+age+dis+rad+tax+ptratio+medv,layout =c(3,3))
```

Marginal Model Plots
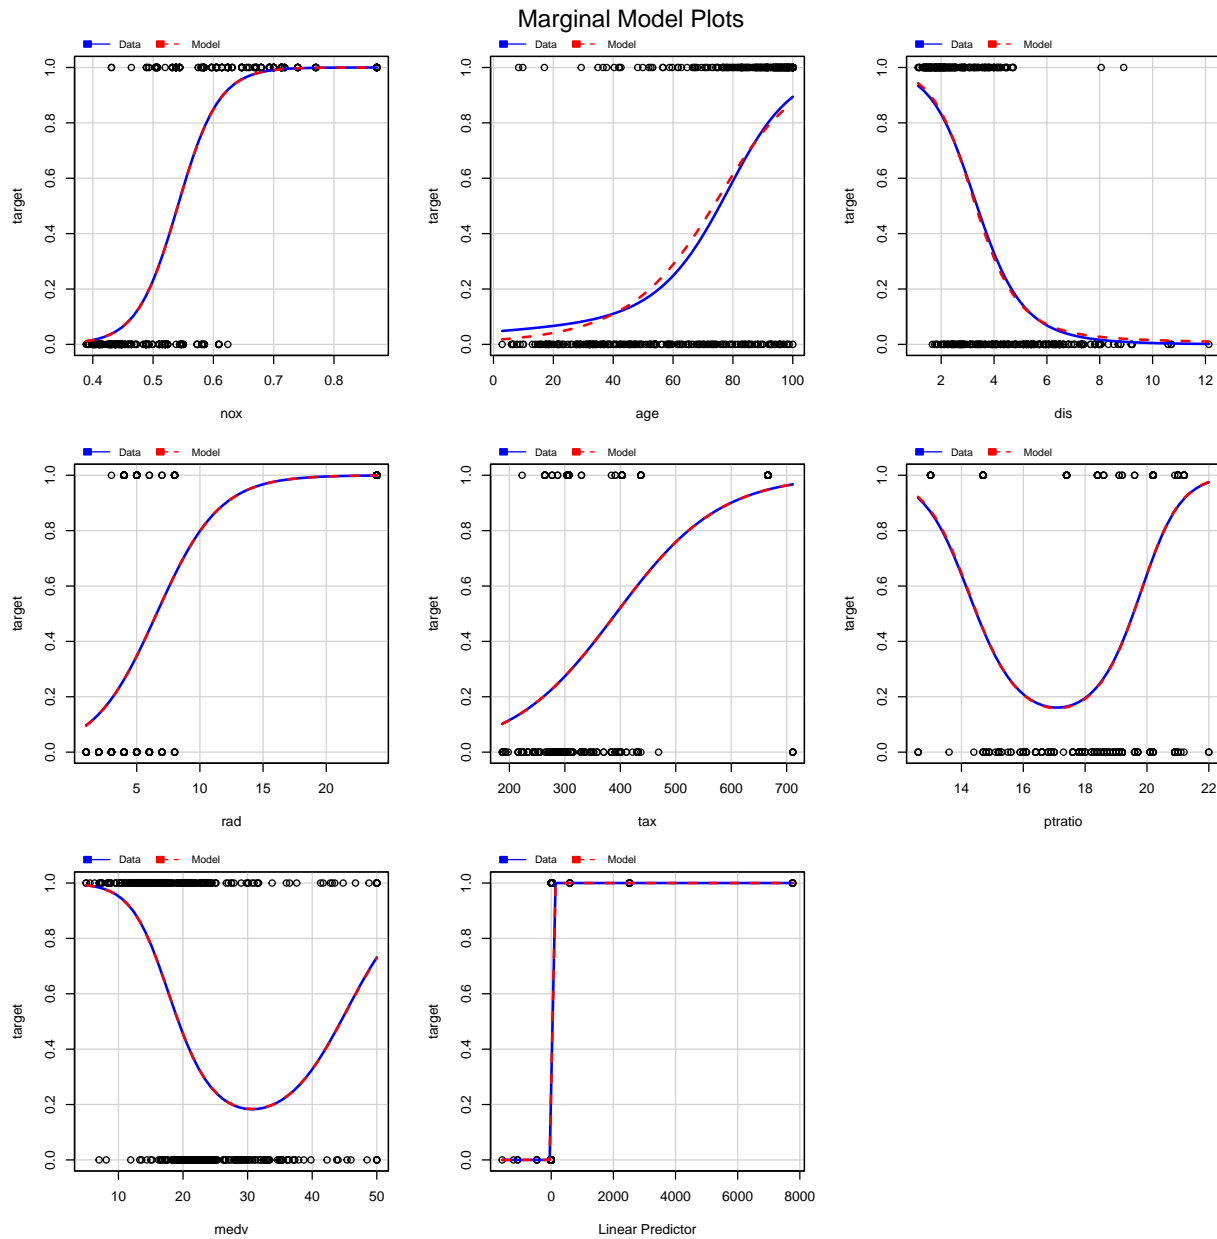
All variables follow nearly the same as the nonparametric estimations. Our model is fitting well to the data.

**Residual Plots**

```
residual_df <- mutate(train_df, residuals=residuals(model_chi_0.1,type="deviance"),
                      linpred=predict(model_chi_0.1,type = "link"))
gdf <- group_by(residual_df, cut(linpred, breaks=unique(quantile(linpred,(1:100)/101))))
diagdf <- summarise(gdf, residuals=mean(residuals), linpred=mean(linpred))
plot(residuals ~ linpred, diagdf, xlab="linear predictor",xlim=c(-20,20))
```
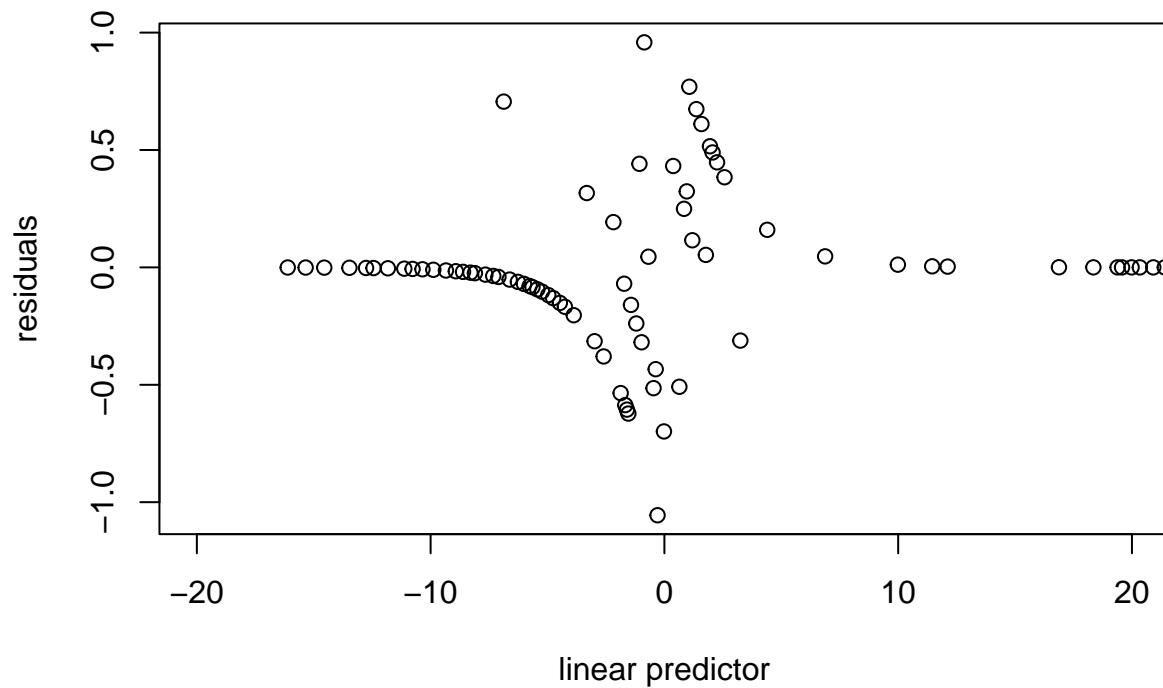
The deviance residual vs linear predictor plot shows that our model is valid. The model is producing accurate predictions at the two ends. The errors around the match point 0 are independent and random.

Let's also check the residual plots with individual predictors

```r
predictors <- c("nox","age","dis","rad","tax","ptratio","medv")

residual_df <- mutate(train_df, residuals=residuals(model_chi_0.1,type="deviance"))
gg_plots <- list()

for (i in c(1:length(predictors))) {
    gdf <- group_by(residual_df, .dots = predictors[i])
    diagdf <- summarise(gdf, residuals=mean(residuals))
    print(ggplot(diagdf, aes_string(x=predictors[i],y="residuals")) + geom_point())
}
```
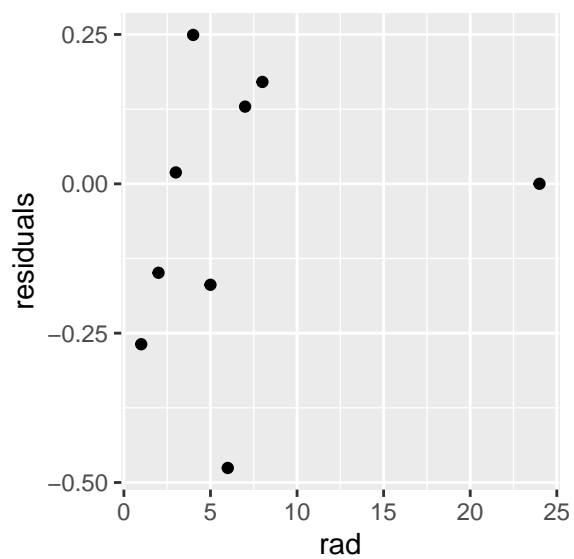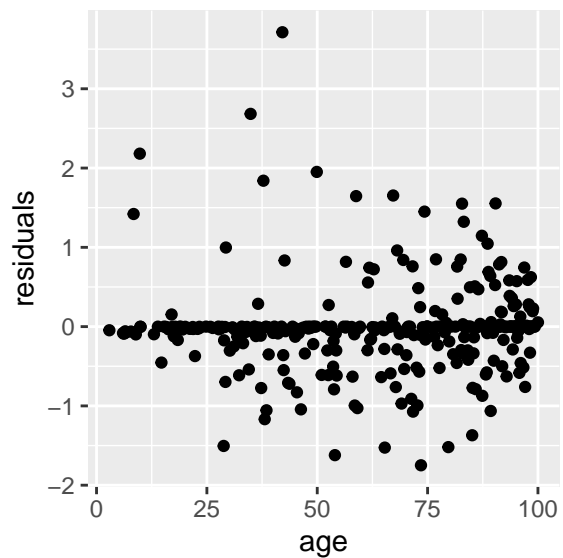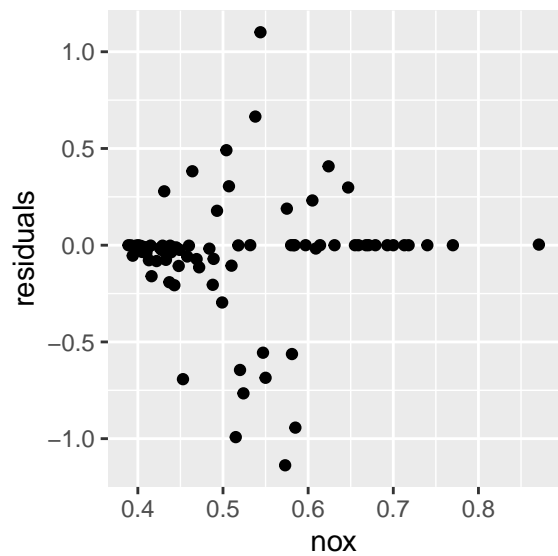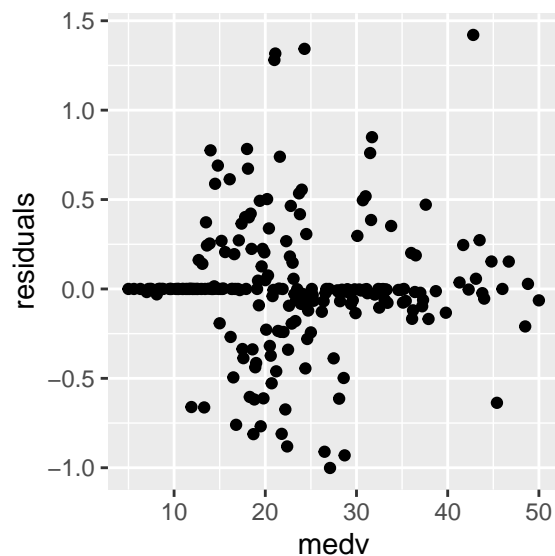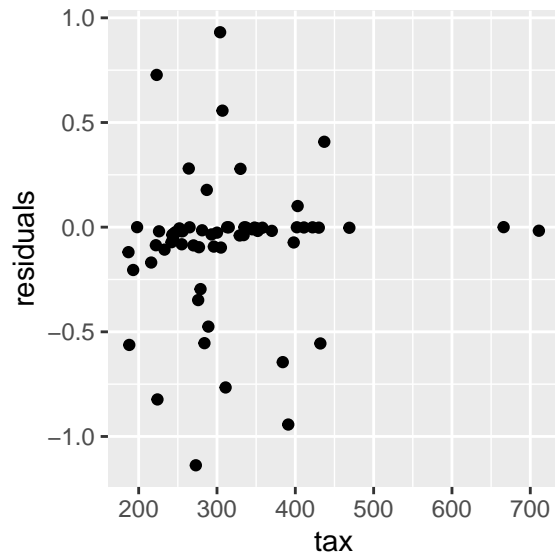
The residuals in each plots are centered at 0, mostly independent and with roughly the same variance, except a few outliers. We conclude that our model does not have notable violation against its validity. The residuals for *nox* and *dis* seems to be heteroscedastic. Given this is a logistic regression with binary data, this phenomena is acceptable.
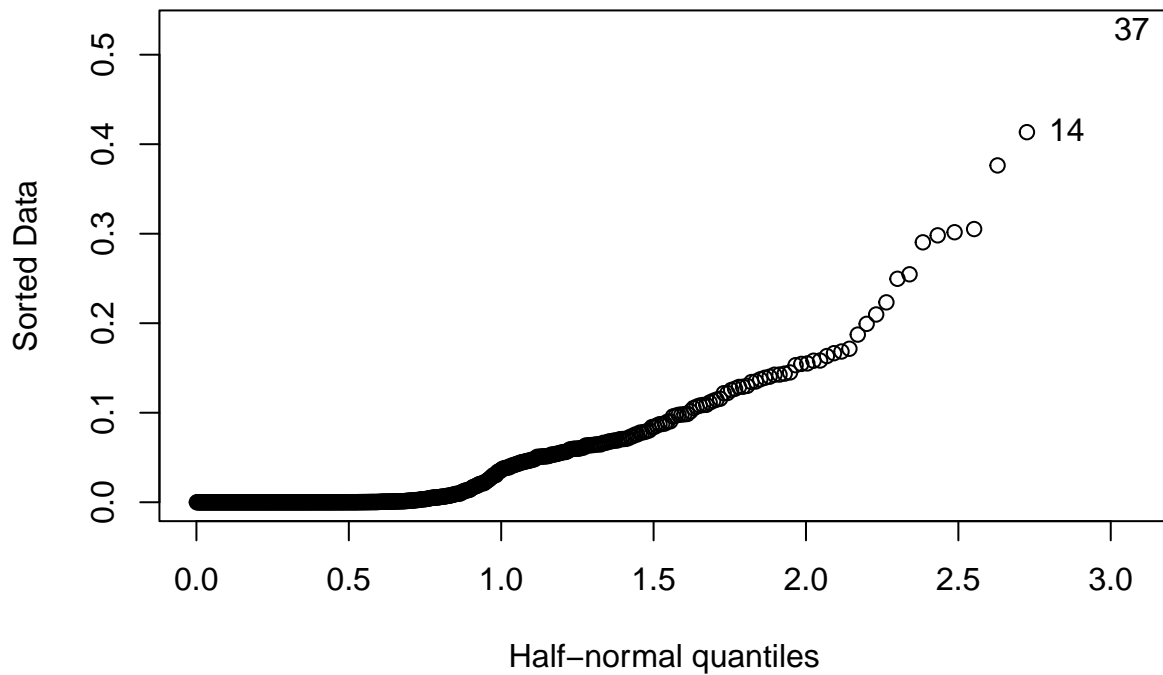
**Q-Q Plot and half normal plot**

```
qqnorm(residuals(model_chi_0.1))
```

## Normal Q–Q Plot



The Q-Q plot seems to be fine given it's a logistic regression with binary data

```
halfnorm(hatvalues(model_chi_0.1))
```

The half normal plots shows case 14 and 37 have high leverage.

By looking at the details of the cases, there is nothing extreme in the values.

```
train_df[c(14,37),]
```

```
##    zn indus chas   nox    rm  age    dis rad tax ptratio lstat medv target zn_y
## 14 22  5.86    0 0.431 8.259  8.4 8.9067   7 330    19.1  3.54 42.8      1    1
## 37  0  2.46    0 0.488 7.831 53.6 3.1992   3 193    17.8  4.45 50.0      0    0
##    log_lstat log_medv indus_squared ptratio_squared tax_chas rad_chas
## 14  1.264127 3.756538       34.3396          364.81        0        0
## 37  1.492904 3.912023        6.0516          316.84        0        0
```

Additionally, the predicted link values are close to 0, which confirmed they are not outliers. We would keep them in our model training.

```
predict(model_chi_0.1,train_df[c(14,37),], type="link")
```

```
##         14         37
## -0.5557560 -0.8629056
```

## * Evaluation data prediction

Finally, let's see how our model will predict using the evaluation data set

```
test_df$zn_y <- 0
test_df$zn_y[test_df$zn>0] <- 1

test_df$indus_squared <- test_df$indus^2
test_df$ptratio_squared <- test_df$ptratio^2

test_df$log_lstat <- log(test_df$lstat)
test_df$log_medv <- log(test_df$medv)

test_df$tax_chas <- test_df$tax * test_df$chas
test_df$rad_chas <- test_df$rad * test_df$chas
```

```
test_df$predicted_class <- ifelse(predict(model_chi_0.1,test_df, type = "response") >0.5,1,0)
```

```
ggplot(test_df, aes(predicted_class)) + geom_bar(fill=c("orange","lightblue"))
```



Both target $= 0$ and target $= 1$ are close to 50%, which is a very plausible outcomes. This is the same as we expected since there are 50% of the cases above the median crime rate and 50% of the cases below the median crime rate.